

Distributed Sampled Convex Problems

David Eckman (dje88) and Calvin Wylie (cjlw278)

Abstract

It is well-known that chance constrained programs are in general intractable because of the non-convexity of the feasible region and the requirement of prior information about the distribution of the uncertain parameters. By solving a sampled convex program with sufficiently many samples of the uncertain parameters, one can find solutions that are feasible for the chance constrained program with high probability. Sampled convex problems are more easily solved because they are characterized by a finite number of constraints and require only that one can obtain i.i.d. samples of the uncertain parameters. However, the required number of samples may still result in a sampled convex program that is expensive to solve exactly. We therefore consider a procedure for obtaining approximate solutions to these sampled convex programs by distributing the constraints across processors and solving smaller subproblems, followed by a consensus on the active constraints. We experiment with a classical portfolio optimization problem and study the trade-offs that arise from this method: including wall clock time, expected violation probability, and feasibility probability for the chance constrained program.

Introduction

Many real-world problems in areas as diverse as finance, energy, and emergency services can be modeled by chance constrained optimization problems [1]. For example, chance constraints can represent an acceptable level of risk in an investment or a fixed level of service that is guaranteed. A typical chance constrained program, CCP_ϵ , is given by

$$\begin{aligned} & \underset{x \in \mathcal{X}}{\text{minimize}} && c^T x \\ & \text{subject to} && \mathbb{P}\{f(x, \delta) \leq 0\} \geq 1 - \epsilon. \end{aligned} \tag{CCP}_\epsilon$$

where $\mathcal{X} \subseteq \mathbb{R}^n$ is compact, δ is an uncertain parameter distributed with probability measure \mathbb{P} over an uncertainty set $\Delta \subseteq \mathbb{R}^d$, and $f : \mathcal{X} \times \Delta \mapsto \mathbb{R}$ is convex in x for any fixed δ . The parameter $0 \leq \epsilon \leq 1$ represents the acceptable probability with which the constraint $f(x, \delta) \leq 0$ can be violated. CCP_ϵ is notoriously intractable because the feasible region is non-convex for general probability measures. Moreover, obtaining a convex approximation relies on knowing the distribution \mathbb{P} (or at least the moments) a priori.

Full source code and data for this paper is available at <https://github.com/calvinwylie/pscp>.

One method for finding “good” solutions which are also feasible for CCP_ϵ with high probability is to take N i.i.d. samples $(\delta_1, \dots, \delta_N)$ and solve the corresponding sampled convex program, SCP_N :

$$\begin{aligned} & \underset{x \in \mathcal{X}}{\text{minimize}} && c^T x \\ & \text{subject to} && f(x, \delta_i) \leq 0 \quad i = 1, \dots, N. \end{aligned} \tag{SCP}_N$$

Notice that in contrast to CCP_ϵ , SCP_N is more tractable since it has a finite number of constraints and a convex feasible region. Note that we also need not explicitly know \mathbb{P} , but instead only need to be able to obtain i.i.d. samples from it.

Let x_N^* denote the (random) solution to SCP_N obtained by taking N samples. For a fixed x_N^* , let $V(x_N^*) := \mathbb{P}\{\delta \in \Delta : f(x_N^*, \delta) > 0\}$ be the corresponding violation probability. If $V(x_N^*) \leq \epsilon$, then x_N^* is feasible for CCP_ϵ .

Sampled convex programs do not provide a guarantee that $V(x_N^*) \leq \epsilon$ for a given sample size N . Instead, Calafiore and Campi [2] proved that one can bound the number of samples needed to ensure that $\mathbb{P}^N\{V(x_N^*) > \epsilon\} \leq \beta$ for $0 < \beta \leq 1$, where \mathbb{P}^N is the product probability measure. That is, the optimal solution to SCP_N is feasible for CCP_ϵ with probability greater than $1 - \beta$. Their bound of $N \geq 1/(\epsilon\beta) - 1$ is simple to calculate, but is quite loose in terms of asymptotics. The bound was further tightened in [3] to choosing N to satisfy

$$\mathbb{P}^N\{V(x_N^*) > \epsilon\} \leq \binom{N}{d} (1 - \epsilon)^{N-d}.$$

and then later by Campi and Garatti in [4] to the solution of a binomial equation:

$$\mathbb{P}^N\{V(x_N^*) > \epsilon\} = \sum_{i=0}^{d-1} \binom{N}{i} \epsilon^i (1 - \epsilon)^{N-i}.$$

Note that all of these bounds are independent of the distribution \mathbb{P} , which means that N can be calculated without prior knowledge of \mathbb{P} . Even for modest ϵ and β , the required N may be large enough that the resulting SCP_N instance is very computationally expensive to solve exactly, especially if the number of decision variables, n , is also large.

One idea for addressing this challenge is to decompose SCP_N into smaller sampled convex programs to be solved in parallel before forming some consensus across the subproblems. Because the subproblems have only a fraction of the constraints, they can potentially be solved much quicker than the full SCP_N problem. However, because each subproblem is only considering a subset of the constraints, there is a loss associated with this simplification: the solution returned will likely violate some of the constraints of SCP_N and therefore is less likely to be a feasible solution for CCP_ϵ . But for some applications, the speedup in terms of time to solution due to parallelism may be worth the decrease in the probability that the solution is CCP_ϵ -feasible.

Parallel Procedure for SCP_N

To motivate our proposed procedure, recall that a constraint $\{x : f(x) \leq 0\}$ is *active* at \bar{x} if $f(\bar{x}) = 0$. Assuming some regularity conditions, a feasible convex program in n dimensional

space will have at most n active constraints at the optimal solution. Therefore the problem of finding an optimal solution to SCP_N can be viewed as the problem of identifying the n active constraints at the optimal solution.

A recent paper by Carlone et al. [5] studies solving large-scale convex programs in a distributed fashion. It assumes a directed graph topology on the communication network connecting a set of processors. Initially, each processor is given a subset of the constraints of the full problem with which it solves a convex subproblem. It then identifies the active constraints at that solution and passes them to other processors according to the network topology. Each processor receives the incoming constraints and resolves a convex subproblem with their original set of constraints plus the new constraints. It was proven that after an almost surely finite number of iterations, all processors will reach a consensus about which constraints are active at the optimal solution to the master problem. Note that the objective in [5] is to recover the exact solution to the convex program, whereas we are interested in finding an approximate solution in far less time.

In practice, the procedure in [5] would take far too long to run to completion, especially since there is no bound on the number of iterations needed for consensus. Instead, we will consider a variation of this procedure which only runs for two iterations and in which the topology of the communication network is a star graph in which all “worker” processors pass active constraints to a “master” processor for a second-stage re-solve. Because there are only two iterations, the solution returned will likely not be the exact solution to SCP_N . However this approximate solution can be calculated in less time, and thus it may be worth a loss in the probability of it being a CCP_ϵ -feasible solution.

Suppose that we have available one master processor and p worker processors. For simplicity, suppose that $N = mp$ for some positive integer $m > n$. Our proposed procedure is as follows:

Parallel SCP_N Procedure

1. Acquire N i.i.d. samples $\delta_1, \dots, \delta_N$.
2. For each worker processor $1 \leq j \leq p$, solve the subproblem

$$\begin{aligned} & \underset{x \in \mathcal{X}}{\text{minimize}} && c^T x \\ & \text{subject to} && f(x, \delta_i) \leq 0 \quad i = m(j-1) + 1, \dots, mj. \end{aligned}$$

Let \tilde{x}_j denote the corresponding optimal solution. Determine the set of active constraints at the optimal solution \tilde{x}_j . Assuming that each sampled convex program is fully supported with probability one (i.e. that there are n active constraints), label the samples corresponding to these constraints $\delta_{(1)}^j, \dots, \delta_{(n)}^j$. Pass the set of active constraints to the master processor.

3. On the master processor, solve the sampled convex program

$$\begin{aligned} & \underset{x \in \mathcal{X}}{\text{minimize}} && c^T x \\ & \text{subject to} && f(x, \delta_{(i)}^j) \leq 0 \quad i = 1, \dots, n \text{ and } j = 1, \dots, p. \end{aligned}$$

Let \tilde{x} denote the corresponding optimal solution.

Changing the number of processors p will impact the time to solution (wall clock time) of the procedure in two conflicting ways, as the sizes of both the first-stage subproblems and the second stage master problem are changed. With more processors, the subproblems are smaller, but the total number of constraints passed to the master program (np) is larger. Also, as the number of constraints in the master problem increases, we would expect $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$ to decrease. The main research question of this study will be: *given these trade-offs, what is the optimal number of processors to use?*

We will also consider another version of this problem. Assume that we are willing to commit to the wall clock time needed to solve SCP_N exactly. But instead of solving SCP_N we want to use the time to take more than N samples and apply the distributed procedure in hopes of increasing the probability that the returned solution is feasible for CCP_ϵ . *What then is the best number of processors and number of constraints per processor?*

Example Problem

We consider the classical portfolio optimization problem of selecting an allocation y for investing one dollar over n assets with uncertain returns r_i for $i = 1, \dots, n$. Of these assets, we assume that the n th asset has a fixed payout; it might, for example, represent a bond with a certain return. Each allocation of assets defines a portfolio. Our objective is to choose a portfolio that maximizes the value-at-risk (VaR), denoted by t , for an acceptable risk level ϵ . Here VaR refers to the ϵ -percentile of the distribution of returns for a particular portfolio as opposed to the (more traditional) ϵ -percentile of the distribution of losses.

This setting gives rise to the following chance constrained program:

$$\begin{aligned}
 & \underset{y, t}{\text{maximize}} && t \\
 & \text{subject to} && \mathbb{P} \left\{ \sum_{i=1}^{n-1} r_i y_i + r_n y_n \geq t \right\} \geq 1 - \epsilon \\
 & && \sum_{i=1}^n y_i = 1 \\
 & && y_i \geq 0 \quad i = 1, \dots, n.
 \end{aligned} \tag{Portfolio CCP}$$

We will assume that the vector of uncertain returns (r_1, \dots, r_{n-1}) is distributed according to a multivariate log-normal distribution with mean vector μ and covariance matrix Σ . A log-normal distribution is a typical model for asset return because it is unimodal with support on the positive real line. A more sophisticated approach might sample the returns from an empirical distribution function built from past data. In the experiments we perform, we take $n = 200$, with one asset fixed at a return of 1.05. The rest of the assets are uncorrelated, with means increasing on the interval $[1.06, 1.23]$ and variances increasing on the interval $[0.05, 1.045]$.

By taking samples $r^{(j)}$ for $j = 1, \dots, N$, the corresponding sampled convex program is

$$\begin{aligned}
& \underset{y, t}{\text{maximize}} && t \\
& \text{subject to} && \sum_{i=1}^{n-1} r_i^{(j)} y_i + r_n^{(j)} y_n \geq t \quad j = 1, \dots, N \\
& && \sum_{i=1}^n y_i = 1 \\
& && y_i \geq 0 \quad i = 1, \dots, n.
\end{aligned} \tag{Portfolio SCP}_N$$

Portfolio SCP_N has several nice properties:

- The objective function and constraints of Portfolio SCP_N are linear. Therefore we can use standard linear programming techniques to solve the sampled convex program and subproblems.
- Because of the fixed asset n , it is possible to easily find a basic feasible solution by taking $y_n = 1$, $y_i = 0$ for $i = 1, \dots, n - 1$ and $t = r_n$.
- The feasible allocations y sit within a unit simplex.

In our experiments, we took $\epsilon = 0.05$ and $\beta = 1 \times 10^{-5}$ and calculated a required sample size of $N = 5312$.

Implementation

To solve our example portfolio optimization problem, we implemented a C interface to solve linear programs in parallel using the GNU Linear Programming Kit (GLPK) ¹. We chose to use the Simplex algorithm to solve our linear programs. As our constraints are very dense, interior point methods would likely be very slow. To provide parallelism, we coded to the Message Passing Interface (MPI) standard. The code was compiled with the Intel C compiler (icc), and experiments performed on a single 12-core Intel Xeon E5-2620 v3 processor.

Given a linear program with m constraints and n variables, the folklore of the simplex method is that it takes $O(m)$ pivots to reach the optimal solution for a “typical” linear program (the worst case is always $2^n - 1$). Thus we should observe a roughly linear speedup when decomposing our sampled program into smaller subproblems. Also given a problem where $m \gg n$, it may make sense to solve the dual program instead (which can be done through the so-called Dual Simplex Algorithm), which has only n constraints.

Figure 1 shows some experimental timings for the simplex and dual simplex methods applied to our problem with $n = 200$ assets on a single core. It shows a superlinear increase in wall clock time with respect to the number of constraints.

However, we do not realize any speedup through solving the dual program. We guess this is because our sampled constraints are on average very similar to each other, and thus not as many pivots are needed to reach an optimal extreme point as would be required in a more “standard” linear program.

¹<https://www.gnu.org/software/glpk/>




Figure 1: Wall Clock Times for Simplex and Dual Simplex Methods.

Performance Analysis

The program spends time in two main areas: communicating the constraint data through MPI and solving the linear programs in GLPK. As these are two mature and heavily optimized libraries, manual performance tuning of the routines called would likely not lead to significant speedups. We can however, reason about the overall performance of our program at a high level.

Assuming that the number of constraints is held constant, as Figure 1 suggests, we should see a linear decrease in the wall clock time of the simplex algorithm when distributing the constraints across different processors. The time required for the final solve on the master processor (using active constraints from the subproblems) should be small, as we can warm-start this optimization with an optimal basis from a subproblem.

However, as we increase the number of processors, the time to communicate the initial constraints (through an `MPI_Scatter`) and the active constraints (through `MPI_Gather`) will increase. Through profiling, we verified that there is a roughly $O(p^2)$ relationship between communication time and the number of processors p .

This suggests that we should initially see decreases in wall clock time as we add processors, but at some point the communication time will start to dominate the savings gained by splitting our problem. Figure 2 shows this experimentally.




Figure 2: Wall Clock Times by Number of Processors.

Experiments

Let $\tilde{x} := (\tilde{y}, \tilde{t})$ denote the solution returned by running the parallel SCP_N procedure. Even though the distribution \mathbb{P} is known, calculating $V(\tilde{x})$ exactly is intractable. Instead we estimated $V(\tilde{x})$ via Monte Carlo simulation, generating $R = 5000$ i.i.d. realizations of the returns r and calculating

$$\hat{V}(\tilde{x}) = \frac{1}{R} \sum_{j=1}^R \mathbf{1} \left\{ \sum_{i=1}^{n-1} r_i^{(j)} \tilde{y}_i + r_n^{(j)} \tilde{y}_n < \tilde{t} \right\}.$$

We performed $M = 400$ macro-replications of this procedure, with corresponding solutions $\tilde{x}^{(k)}$ for $k = 1, \dots, M$, and calculated the following point estimates of $\mathbb{E}[V(\tilde{x})]$ and $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$:

$$\widehat{\mathbb{E}[V(\tilde{x})]} := \frac{1}{M} \sum_{k=1}^M \hat{V}(\tilde{x}^{(k)}),$$

and

$$\mathbb{P}^N\{\widehat{V(\tilde{x})} > \epsilon\} := \frac{1}{M} \sum_{k=1}^M \mathbf{1} \left\{ \hat{V}(\tilde{x}^{(k)}) > \epsilon \right\}.$$

We used the Clopper-Pearson exact binomial test to form approximate confidence intervals for these two performance measures.

When we divided the 5312 constraints evenly among p processors, our estimates for $\mathbb{E}[V(\tilde{x})]$ and $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$ are shown in Figures 3 and 4.




Figure 3: $\mathbb{E}[V(\tilde{x})]$ by Number of Processors.

Though at four processors and above the violation probability $V(\tilde{x})$ is less than ϵ in expectation, the probability of CCP_ϵ constraint violation $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$ is still very high (recall that we are aiming for $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\} \leq \beta = 10^{-5}$).

Figure 5 shows the trade-off between wall clock time and $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$. This Figure shows that despite returning solutions with high probabilities of infeasibility, the parallel SCP_N procedure could run in a fraction of the time of the serial procedure. This suggests that we could use some of the time savings to run the parallel procedure on more than N samples, recovering solutions with feasibility probabilities competitive with the serial procedure, with still significant speedup.

In another set of experiments, we increased the number of samples given to each worker processor, typically doubling or quadrupling. For example, with $N = 5312$, running the parallel procedure with eight processors means that each processor receives $N/8 = 664$ samples. Here we ran the same procedure on eight processors, except with 1328 and 2656 constraints for each processor (so now the total number of samples generated was 10624 and 21248).

Figure 6 shows the wall clock time results for the instances we tested. Observe that the lowest point on each colored curve corresponds to the case of running the original problem of $N = 5312$ constraints. From Figure 6, we can see the increases in wall clock time for a fixed number of processors is again superlinear. In addition, doubling the number of constraints

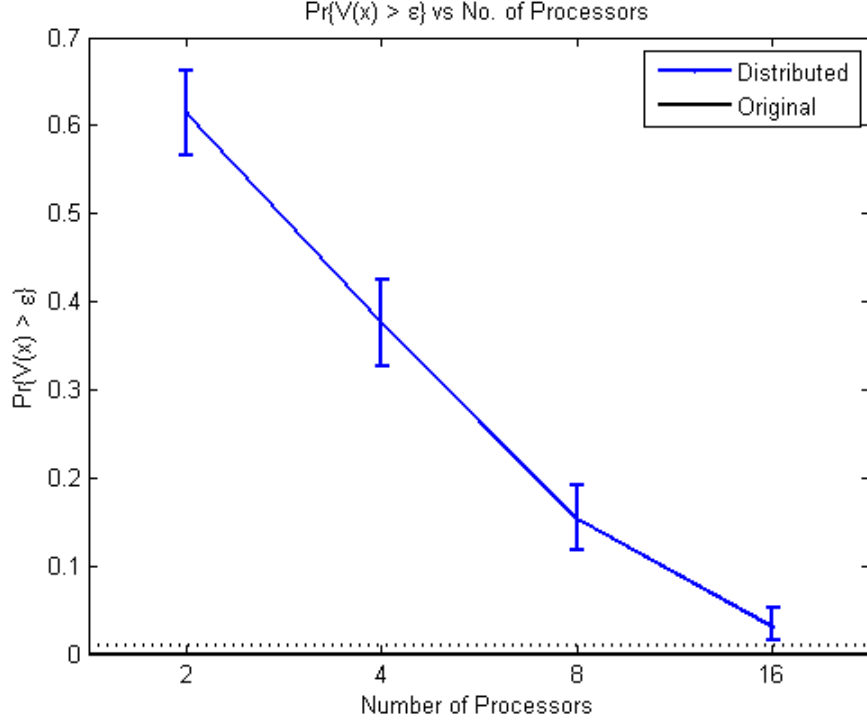


Figure 4: $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$ by Number of Processors.

per processor (for $p = 4, 8$, and 16) appears to outperform the serial procedure in terms of wall clock time and quadrupling the number of constraints per processor (for $p = 4$ and 8) has comparable times to the serial procedure.

Figures 7 and 8 show $\mathbb{E}[V(\tilde{x})]$ and $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$, respectively, for these same instances. On these plots, the highest point on each colored curve corresponds to the case of running the original problem. As expected, both performance measures improve as the number of constraints per processor are increased. It appears that doubling the number of constraints per processor is enough for the parallel procedure to match or surpass the performance of the serial procedure.

Finally, the wall clock time and $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$ are shown together in Figure 9. Each colored line represents a frontier for a fixed number of processors where the left-point on each curve is the original problem. The number of constraints per processor increases as one moves to the right along each curve. Based on this plot, an “optimal” number of configuration appears to be eight processors with 1328 constraints per processor (the middle red point) because it attains the $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$ of the serial solution in less than half the time.

Conclusions

We presented a parallel procedure for obtaining approximate solutions to sampled convex programs. For a modest-sized portfolio optimization problem, we observed that this parallel procedure can dramatically reduce the run time with some loss in CCP_ϵ feasibility. In our




Figure 5: Trade-off between Time and $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$.

MPI implementation we observed marginally decreasing gains in run-time performance as the number of processors increased, due to increased communication and the need to solve a larger second-stage problem. We believe that for more complex sampled convex programs—those with more decision variables and constraints or those of harder complexity than linear programs—the benefits of a distributed approach would be even more pronounced. When we experimented with adjusting the number of constraints per processor, our main conclusion was that one can typically sample twice as much while achieving faster run-times and CCP_ϵ -feasibility competitive with the serial procedure.

Open Extensions

If we had more time, there are a number of open areas we would have wanted to explore:

- Experiment on an actual financial data set with forecasts for asset returns.
- Compare the quality of the solution returned by our procedure with other solutions found by solving over convex inner approximations of the CCP_ϵ feasible region.
- Extend our procedure to run multiple iterations of the worker/master sequence in hopes of better recovering the optimal active constraints.
- Test for large problem sizes and other problems whose solution methods are of harder complexity than those of linear programming.

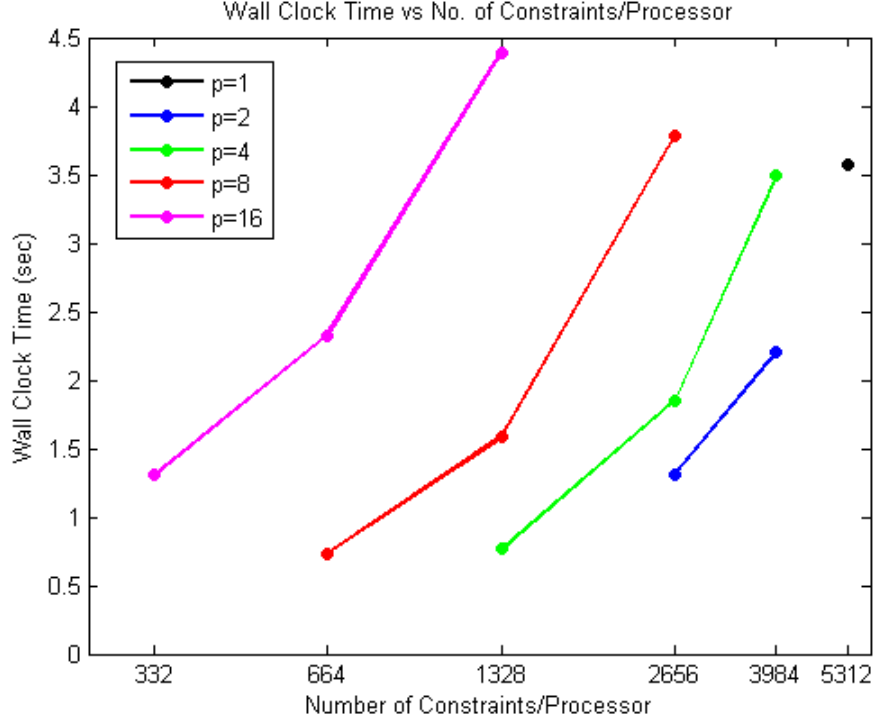


Figure 6: Wall Clock Times by Number of Constraints/Processor.

References

- [1] A. BEN-TAL, L. EL GHAOU, AND A. NEMIROVSKI, *Robust optimization*, Princeton University Press, 2009.
- [2] G. CALAFIORE AND M. C. CAMPI, *Uncertain convex programs: randomized solutions and confidence levels*, Mathematical Programming, 102 (2005), pp. 25–46.
- [3] —, *The scenario approach to robust control design*, IEEE Transactions on Automatic Control, 51 (2006), pp. 742–753.
- [4] M. C. CAMPI AND S. GARATTI, *The exact feasibility of randomized solutions of uncertain convex programs*, SIAM Journal on Optimization, 19 (2008), pp. 1211–1230.
- [5] L. CARLONE, V. SRIVASTAVA, F. BULLO, AND G. C. CALAFIORE, *Distributed random convex programming via constraints consensus*, SIAM Journal on Control and Optimization, 52 (2014), pp. 629–662.




Figure 7: $E[V(\tilde{x})]$ by Number of Constraints/Processor.

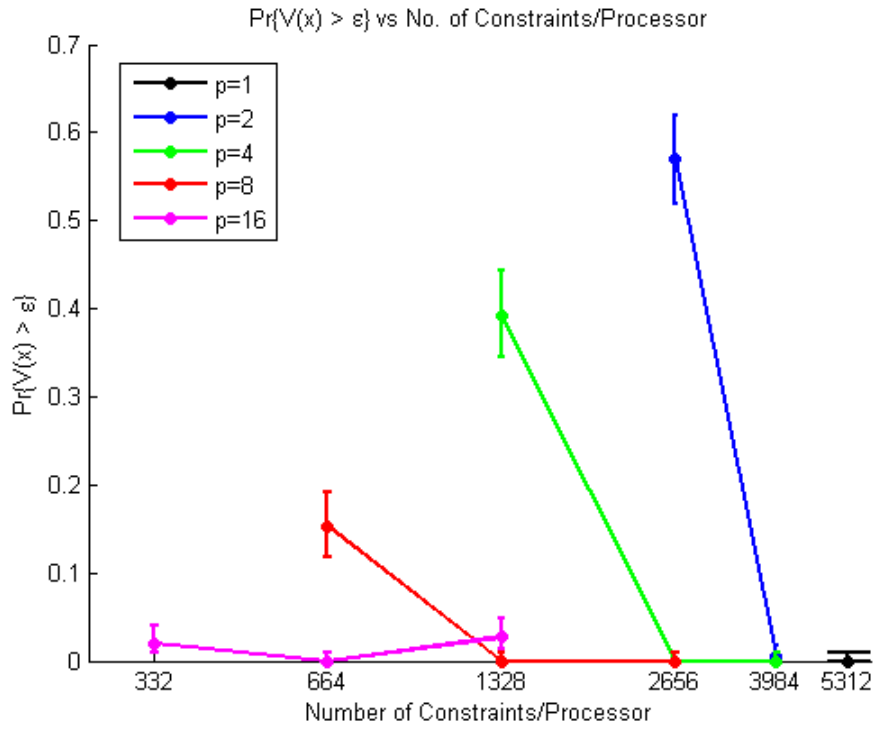


Figure 8: $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$ by Number of Constraints/Processor.




Figure 9: Trade-off between Time and $\mathbb{P}^N\{V(\tilde{x}) > \epsilon\}$